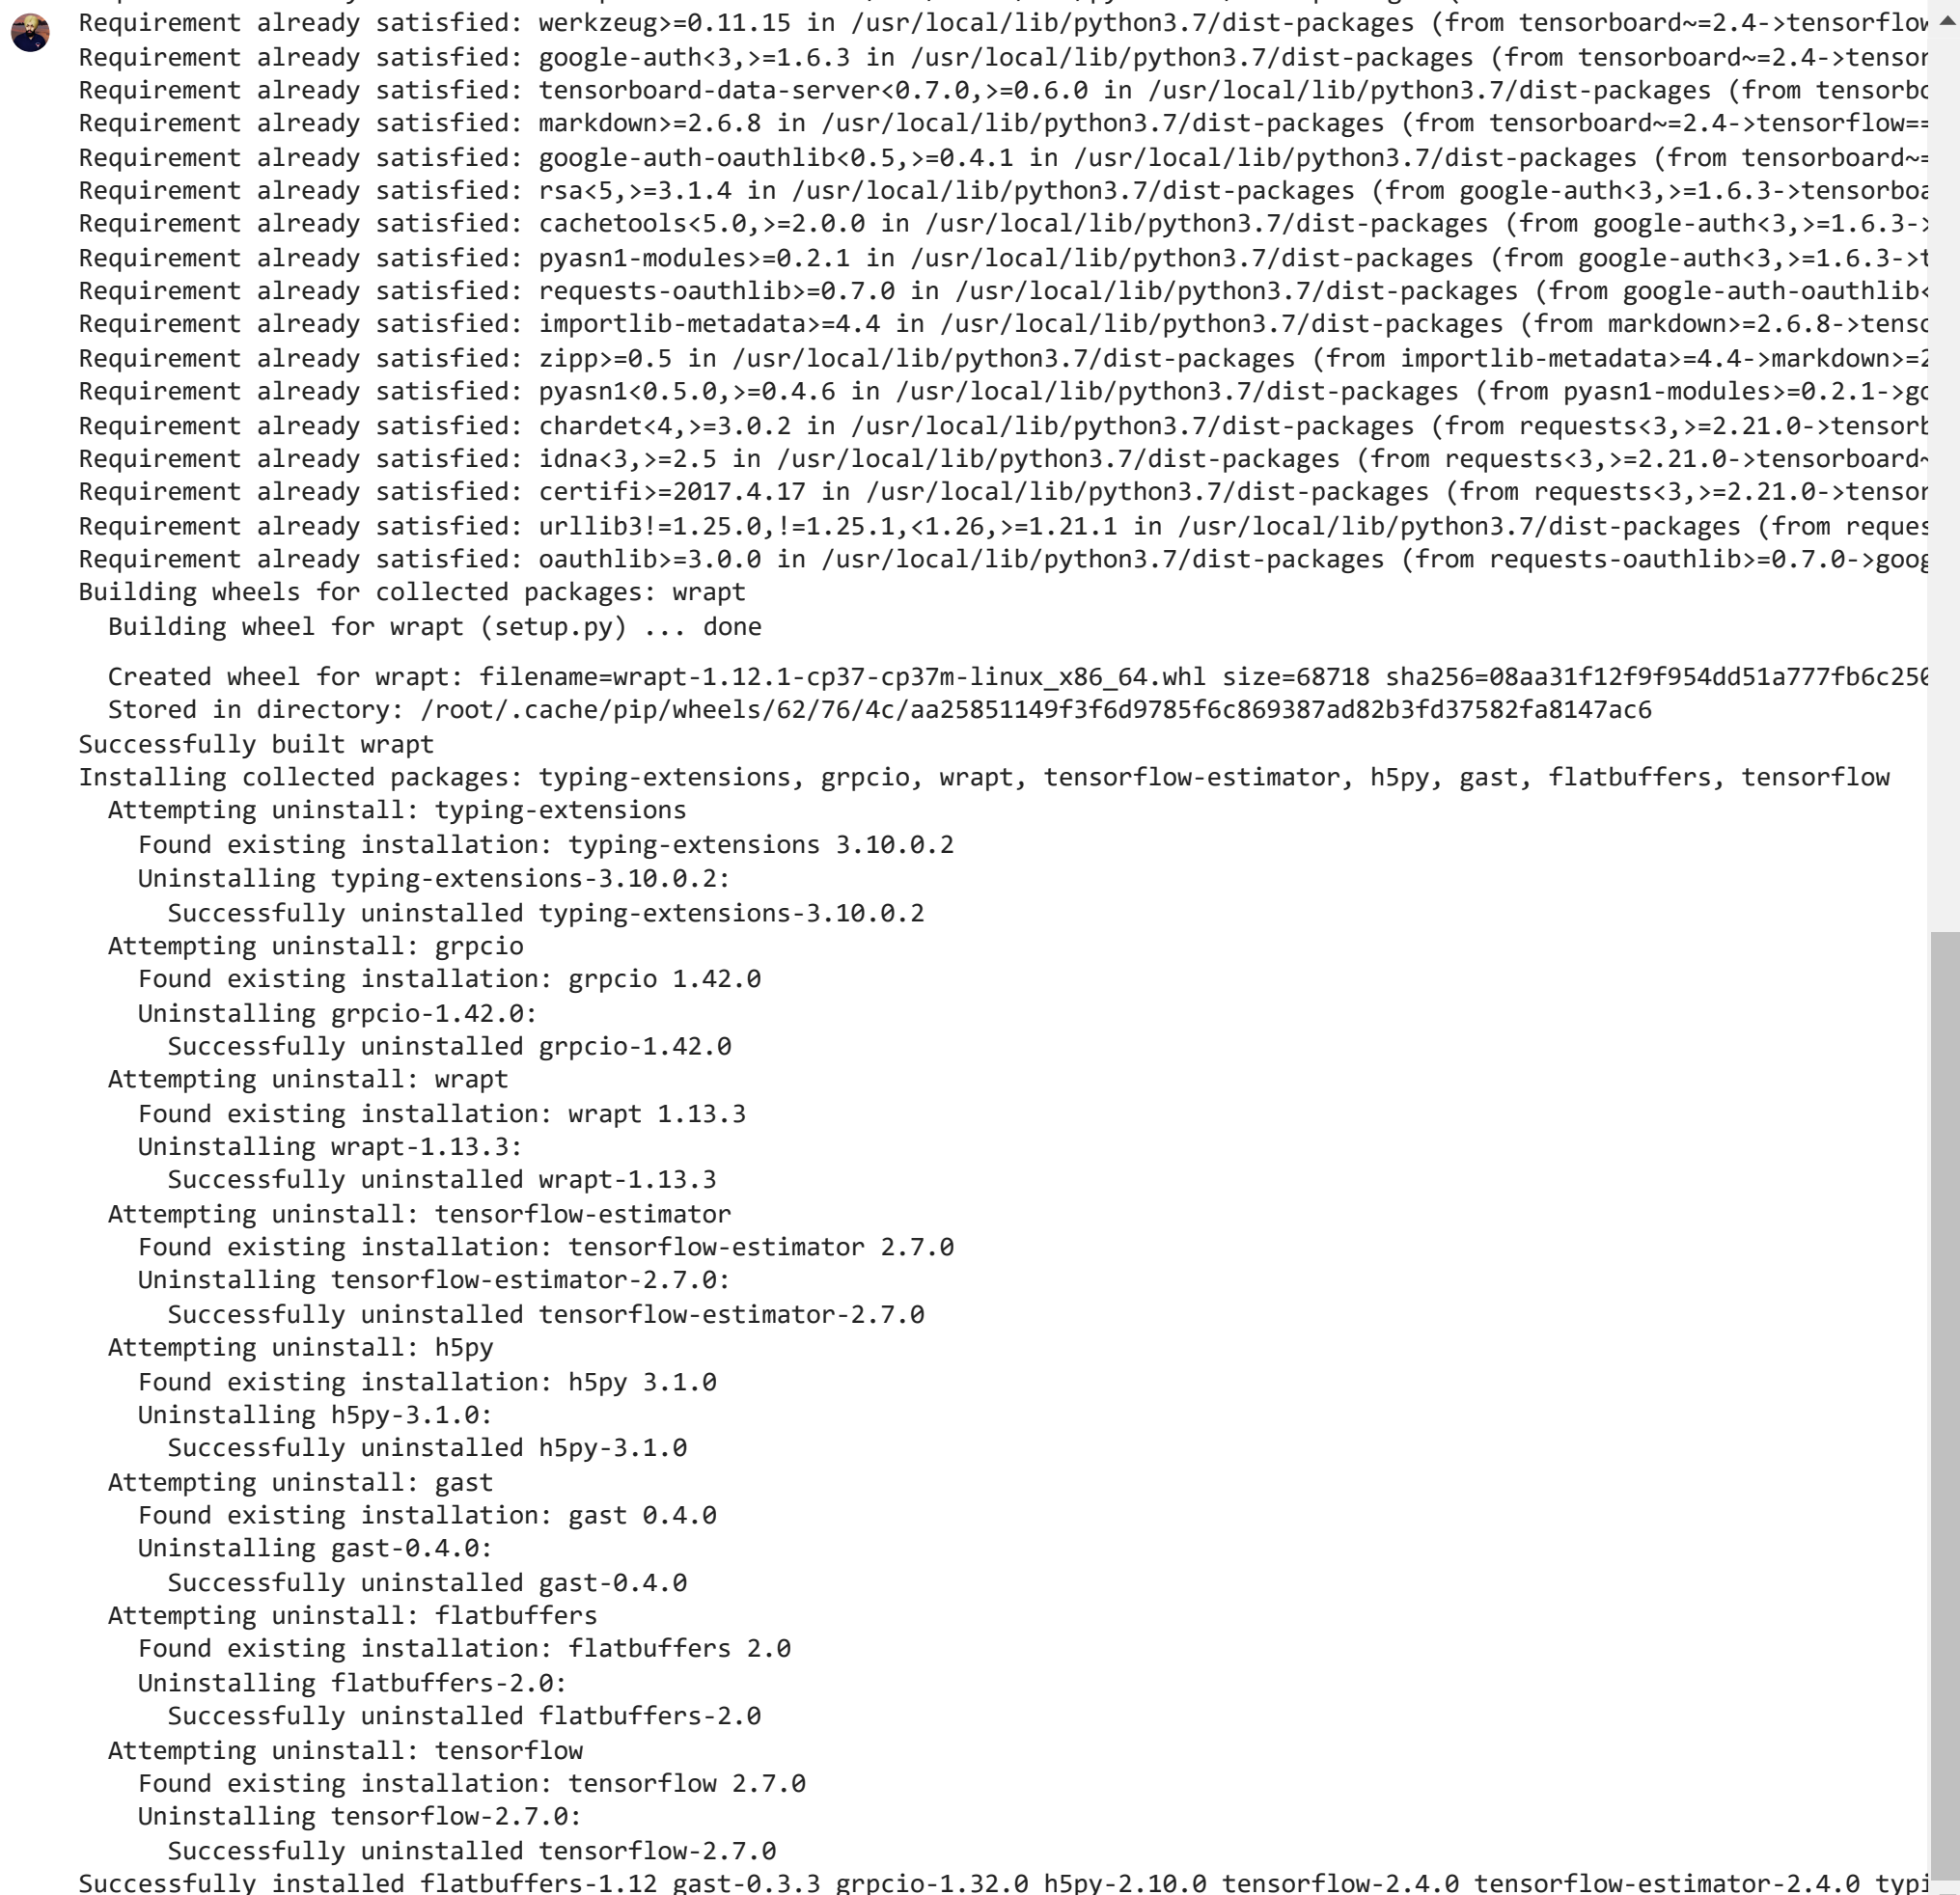


▼ Download and Visualize Data

```
# import dataset
```

```
!wget https://github.com/hfg-gmuend/openmoji/releases/latest/download/openmoji-72x72-color.zip
!mkdir emojis
!unzip -q openmoji-72x72-color.zip -d ./emojis
!pip install tensorflow==2.4
```



```
Requirement already satisfied: werkzeug>=0.11.15 in /usr/local/lib/python3.7/dist-packages (from tensorboard~=2.4->tensorflow)
Requirement already satisfied: google-auth<3,>=1.6.3 in /usr/local/lib/python3.7/dist-packages (from tensorboard~=2.4->tensorflow)
Requirement already satisfied: tensorboard-data-server<0.7.0,>=0.6.0 in /usr/local/lib/python3.7/dist-packages (from tensorboard~=2.4->tensorflow)
Requirement already satisfied: markdown>=2.6.8 in /usr/local/lib/python3.7/dist-packages (from tensorboard~=2.4->tensorflow)
Requirement already satisfied: google-auth-oauthlib<0.5,>=0.4.1 in /usr/local/lib/python3.7/dist-packages (from tensorboard~=2.4->tensorflow)
Requirement already satisfied: rsa<5,>=3.1.4 in /usr/local/lib/python3.7/dist-packages (from google-auth<3,>=1.6.3->tensorboard)
Requirement already satisfied: cachetools<5.0,>=2.0.0 in /usr/local/lib/python3.7/dist-packages (from google-auth<3,>=1.6.3->tensorboard)
Requirement already satisfied: pyasn1-modules>=0.2.1 in /usr/local/lib/python3.7/dist-packages (from google-auth<3,>=1.6.3->tensorboard)
Requirement already satisfied: requests-oauthlib>=0.7.0 in /usr/local/lib/python3.7/dist-packages (from google-auth-oauthlib<0.5,>=0.4.1->tensorboard)
Requirement already satisfied: importlib-metadata>=4.4 in /usr/local/lib/python3.7/dist-packages (from markdown>=2.6.8->tensorflow)
Requirement already satisfied: zipp>=0.5 in /usr/local/lib/python3.7/dist-packages (from importlib-metadata>=4.4->markdown>=2.6.8->tensorflow)
Requirement already satisfied: pyasn1<0.5.0,>=0.4.6 in /usr/local/lib/python3.7/dist-packages (from pyasn1-modules>=0.2.1->google-auth<3,>=1.6.3->tensorboard)
Requirement already satisfied: chardet<4,>=3.0.2 in /usr/local/lib/python3.7/dist-packages (from requests<3,>=2.21.0->tensorboard)
Requirement already satisfied: idna<3,>=2.5 in /usr/local/lib/python3.7/dist-packages (from requests<3,>=2.21.0->tensorboard)
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.7/dist-packages (from requests<3,>=2.21.0->tensorboard)
Requirement already satisfied: urllib3!=1.25.0,!1.25.1,<1.26,>=1.21.1 in /usr/local/lib/python3.7/dist-packages (from requests<3,>=2.21.0->tensorboard)
Requirement already satisfied: oauthlib>=3.0.0 in /usr/local/lib/python3.7/dist-packages (from requests-oauthlib>=0.7.0->google-auth-oauthlib<0.5,>=0.4.1->tensorboard)
Building wheels for collected packages: wrapt
  Building wheel for wrapt (setup.py) ... done
  Created wheel for wrapt: filename=wrapt-1.12.1-cp37-cp37m-linux_x86_64.whl size=68718 sha256=08aa31f12f9f954dd51a777fb6c256
  Stored in directory: /root/.cache/pip/wheels/62/76/4c/aa25851149f3f6d9785f6c869387ad82b3fd37582fa8147ac6
Successfully built wrapt
Installing collected packages: typing-extensions, grpcio, wrapt, tensorflow-estimator, h5py, gast, flatbuffers, tensorflow
Attempting uninstall: typing-extensions
  Found existing installation: typing-extensions 3.10.0.2
  Uninstalling typing-extensions-3.10.0.2:
    Successfully uninstalled typing-extensions-3.10.0.2
Attempting uninstall: grpcio
  Found existing installation: grpcio 1.42.0
  Uninstalling grpcio-1.42.0:
    Successfully uninstalled grpcio-1.42.0
Attempting uninstall: wrapt
  Found existing installation: wrapt 1.13.3
  Uninstalling wrapt-1.13.3:
    Successfully uninstalled wrapt-1.13.3
Attempting uninstall: tensorflow-estimator
  Found existing installation: tensorflow-estimator 2.7.0
  Uninstalling tensorflow-estimator-2.7.0:
    Successfully uninstalled tensorflow-estimator-2.7.0
Attempting uninstall: h5py
  Found existing installation: h5py 3.1.0
  Uninstalling h5py-3.1.0:
    Successfully uninstalled h5py-3.1.0
Attempting uninstall: gast
  Found existing installation: gast 0.4.0
  Uninstalling gast-0.4.0:
    Successfully uninstalled gast-0.4.0
Attempting uninstall: flatbuffers
  Found existing installation: flatbuffers 2.0
  Uninstalling flatbuffers-2.0:
    Successfully uninstalled flatbuffers-2.0
Attempting uninstall: tensorflow
  Found existing installation: tensorflow 2.7.0
  Uninstalling tensorflow-2.7.0:
    Successfully uninstalled tensorflow-2.7.0
Successfully installed flatbuffers-1.12 gast-0.3.3 grpcio-1.32.0 h5py-2.10.0 tensorflow-2.4.0 tensorflow-estimator-2.4.0 typi
```

```
# import libraries
%matplotlib inline
```

```
import tensorflow as tf
import numpy as np
import matplotlib.pyplot as plt
import os
```

```
from PIL import Image, ImageDraw
from tensorflow.keras.layers import Input, Dense, Flatten, Conv2D, MaxPool2D, BatchNormalization, Dropout
```

```

print('Check if we are using TensorFlow 2.4')
print('Using TensorFlow version', tf.__version__)


    Check if we are using TensorFlow 2.4
    .. _ _ _ _ _

# emojis actually using in this project
emojis = {
    0: {'name': '8ball', 'file': '1F3B1.png'},
    1: {'name': 'tennisball', 'file': 'tennisball.png'},
    2: {'name': 'basketball', 'file': '1F3C0.png'},
    3: {'name': 'football', 'file': '26BD.png'},
    4: {'name': 'baseball', 'file': '26BE.png'},
    5: {'name': 'leatherball', 'file': '192950.png'}
}

# place images in larger images - help synthesis data for localization
plt.figure(figsize=(9, 9))

for i, (j, e) in enumerate(emojis.items()):
    plt.subplot(3, 3, i + 1)
    plt.imshow(plt.imread(os.path.join('emojis', e['file'])))
    plt.xlabel(e['name'])
    plt.xticks([])
    plt.yticks([])
plt.show()

```



▼ Create Examples

```

# loads emojis and assign a key - image for each class in emojis dictionary
for class_id, values in emojis.items():
    png_file = Image.open(os.path.join('emojis', values['file'])).convert('RGBA')
    png_file.load()
    new_file = Image.new("RGB", png_file.size, (255, 255, 255))
    new_file.paste(png_file, mask=png_file.split()[3])
    emojis[class_id]['image'] = new_file

```

```

emojis

{0: {'file': '1F3B1.png',
     'image': <PIL.Image.Image image mode=RGB size=72x72 at 0x7F96EE4CF190>,
     'name': '8ball'},
 1: {'file': 'tennisball.png',
     'image': <PIL.Image.Image image mode=RGB size=72x72 at 0x7F96EE4CF090>,
     'name': 'tennisball'},
 2: {'file': '1F3C0.png',
     'image': <PIL.Image.Image image mode=RGB size=72x72 at 0x7F96EE4CF4D0>,
     'name': 'basketball'},
 3: {'file': '26BD.png',
     'image': <PIL.Image.Image image mode=RGB size=72x72 at 0x7F96EE4CFD50>,
     'name': 'football'},
 4: {'file': '26BE.png',
     'image': <PIL.Image.Image image mode=RGB size=72x72 at 0x7F96EE4CF610>,
     'name': 'baseball'},
 5: {'file': '192950.png',
     'image': <PIL.Image.Image image mode=RGB size=72x72 at 0x7F96EE4CFDD0>,
     'name': 'leatherball'}}

```

```
def create_example():
```

```


class_id = np.random.randint(0,6) # randomly choose a emoji
image = np.ones((144,144,3)) * 255 # create a white image
row = np.random.randint(0,72) # place image randomly
col = np.random.randint(0,72) #place image randomly
# place emoji in blank emoji
image[row: row+72, col: col+72, :]= np.array(emojis[class_id]['image'])
# return synthesize image
return image.astype('uint8'), class_id, (row+5)/144 , (col+5)/144      # +10 becuae there is a white space of around 10 pixels in
# /144 is normalization of image

```

```

image, class_id, row, col = create_example()
plt.imshow(image)

```



▼ Plot Bounding Boxes

```

#plotting bounding boxes
def plot_bounding_box(image, gt_coords, pred_coords=[], norm=False): # pass image, ground truth row col coordinates, predicted coordi
    if norm:
        image *=255 # if norm is true, we will denormalize it
        image=image.astype('uint8')
        image = Image.fromarray(image) #to convert image array to pil image
        draw = ImageDraw.Draw(image)

        # extrating row and col from groud truth values
        row, col = gt_coords

        # denormalizing coords
        row *= 144
        col *= 144
        draw.rectangle((col, row, col+62, row+62), outline = 'green', width=3) #+52 becuae iage is of 72 pixels becuae image has buffer

        # now same for pred coords
        if len(pred_coords)==2:
            # extrating row and col from groud truth values
            row, col = pred_coords

            # denormalizing coords
            row *= 144
            col *= 144
            draw.rectangle((col, row, col+62, row+62), outline = 'red', width=3)

    return image

image = plot_bounding_box(image, gt_coords = [row, col])
plt.imshow(image)
plt.title(emojis[class_id]['name'])
plt.show()

```



▼ Data Generator




```
# create endless stream of these randomly generated eg which we will use in our model
def data_generator(batch_size=16):
    #run in endless loop and create example and labels of batch size
    while True:
        x_batch = np.zeros((batch_size, 144, 144, 3)) # 144 = size of image
        y_batch = np.zeros((batch_size, 6)) #9= no of class ids
        bbox_batch = np.zeros((batch_size, 2)) #2 = for row and col values

        # create examples of no of batch size
        for i in range(0, batch_size):
            image, class_id, row, col = create_example()
            x_batch[i] = image/255 # normalize image and 255 because they are pixel values
            y_batch[i, class_id] = 1.0
            bbox_batch[i] = np.array([row,col])
        yield {'image':x_batch},{'class_out': y_batch, 'box_out': bbox_batch}
```

```
example, label = next(data_generator(1))
image = example['image'][0]
class_id = np.argmax(label['class_out'][0]) # to get the actual class id
coords = label['box_out'][0]
```

```
image = plot_bounding_box(image, coords, norm=True)
plt.imshow(image)
plt.title(emojis[class_id]['name'])
plt.show()
```



▼ Model

```
# CNN model
input_ = Input(shape=(144,144,3), name='image')

x = input_

# we will have toatal 5 convolutional blocks
for i in range(0,5):
    n_filters = 2**(4+i)
    x = Conv2D(n_filters, 3, activation='relu')(x) # (x) - input is x
    x = BatchNormalization()(x)
    x = MaxPool2D(2)(x) #pool size od 2x2

x = Flatten()(x)
x = Dense(256, activation = 'relu')(x)

# now connect fully conected layer to our 2 output
class_out = Dense(6, activation='softmax', name = 'class_out')(x) # 9= outputs as we have 9 clases# for classifaction out, we use
box_out = Dense(2, name = 'box_out')(x) # we dont specify any activation as it is regression output and it is linear by default

# now construct the model
model = tf.keras.models.Model(input_, [class_out, box_out])
model.summary()
```

Model: "model_6"

Layer (type)	Output Shape	Param #	Connected to
image (InputLayer)	[(None, 144, 144, 3)]	0	
conv2d_30 (Conv2D)	(None, 142, 142, 16)	448	image[0][0]
batch_normalization_30 (Batch Normalization)	(None, 142, 142, 16)	64	conv2d_30[0][0]
max_pooling2d_30 (MaxPooling2D)	(None, 71, 71, 16)	0	batch_normalization_30[0][0]
conv2d_31 (Conv2D)	(None, 69, 69, 32)	4640	max_pooling2d_30[0][0]
batch_normalization_31 (Batch Normalization)	(None, 69, 69, 32)	128	conv2d_31[0][0]
max_pooling2d_31 (MaxPooling2D)	(None, 34, 34, 32)	0	batch_normalization_31[0][0]
conv2d_32 (Conv2D)	(None, 32, 32, 64)	18496	max_pooling2d_31[0][0]
batch_normalization_32 (Batch Normalization)	(None, 32, 32, 64)	256	conv2d_32[0][0]
max_pooling2d_32 (MaxPooling2D)	(None, 16, 16, 64)	0	batch_normalization_32[0][0]
conv2d_33 (Conv2D)	(None, 14, 14, 128)	73856	max_pooling2d_32[0][0]
batch_normalization_33 (Batch Normalization)	(None, 14, 14, 128)	512	conv2d_33[0][0]
max_pooling2d_33 (MaxPooling2D)	(None, 7, 7, 128)	0	batch_normalization_33[0][0]
conv2d_34 (Conv2D)	(None, 5, 5, 256)	295168	max_pooling2d_33[0][0]
batch_normalization_34 (Batch Normalization)	(None, 5, 5, 256)	1024	conv2d_34[0][0]
max_pooling2d_34 (MaxPooling2D)	(None, 2, 2, 256)	0	batch_normalization_34[0][0]
flatten_6 (Flatten)	(None, 1024)	0	max_pooling2d_34[0][0]
dense_6 (Dense)	(None, 256)	262400	flatten_6[0][0]
class_out (Dense)	(None, 6)	1542	dense_6[0][0]
box_out (Dense)	(None, 2)	514	dense_6[0][0]
Total params: 659,048			
Trainable params: 658,056			
Non-trainable params: 992			

Custom Metric: IoU

```
# intersection over union is the evaluation metric
# to measure the performance of the model - common in finding accuracy in object detector and object localizers
# iou is area of overlap [intesection of 2 boxes] between the predicted bounding box and actual values and
# combining the areas of both minus intersection will give area of union
# divide area of overlap by area of union - IoU values, if 1 - prediction is accurate
```

```
class IoU(tf.keras.metrics.Metric):
    def __init__(self, **kwargs):
        super(IoU, self).__init__(**kwargs)

        self.iou = self.add_weight(name='iou', initializer='zeros')
        self.total_iou = self.add_weight(name='total_iou', initializer='zeros')
        self.num_ex = self.add_weight(name='num_ex', initializer='zeros')

    def update_state(self, y_true, y_pred, sample_weight=None):
        def get_box(y):
            rows, cols = y[:, 0], y[:, 1]
            rows, cols = rows * 144, cols * 144
            y1, y2 = rows, rows + 62
            x1, x2 = cols, cols + 62
            return x1, y1, x2, y2

        def get_area(x1, y1, x2, y2):
            return tf.math.abs(x2 - x1) * tf.math.abs(y2 - y1)

        gt_x1, gt_y1, gt_x2, gt_y2 = get_box(y_true)
        p_x1, p_y1, p_x2, p_y2 = get_box(y_pred)
```

```

    i_x1 = tf.maximum(gt_x1, p_x1)
    i_y1 = tf.maximum(gt_y1, p_y1)
    i_x2 = tf.minimum(gt_x2, p_x2)
    i_y2 = tf.minimum(gt_y2, p_y2)

    i_area = get_area(i_x1, i_y1, i_x2, i_y2)
    u_area = get_area(gt_x1, gt_y1, gt_x2, gt_y2) + get_area(p_x1, p_y1, p_x2, p_y2) - i_area

    iou = tf.math.divide(i_area, u_area)
    self.num_ex.assign_add(1)
    self.total_iou.assign_add(tf.reduce_mean(iou))
    self.iou = tf.math.divide(self.total_iou, self.num_ex)

def result(self):
    return self.iou

def reset_state(self):
    self.iou = self.add_weight(name='iou', initializer='zeros')
    self.total_iou = self.add_weight(name='total_iou', initializer='zeros')
    self.num_ex = self.add_weight(name='num_ex', initializer='zeros')

```

▼ Task 8: Compile the Model

```

model.compile(
    # specify the loss for fiffereent outputs
    loss={
        'class_out': 'categorical_crossentropy', # or classifaction output
        'box_out': 'mse' #for regression output
    },
    optimizer = tf.keras.optimizers.Adam(learning_rate=1e-3),

    #set differnet metric for different output
    metrics={
        'class_out': 'accuracy',
        'box_out': IoU(name='iou')
    }
)

```

▼ Custom Callback: Model Testing

```

def test_model(model, test_datagen):
    example, label = next(test_datagen)
    x = example['image']
    y = label['class_out']
    box = label['box_out']

    pred_y, pred_box = model.predict(x)

    pred_coords = pred_box[0] # [0] beacuse we want only one example
    gt_coords = box[0]
    pred_class = np.argmax(pred_y[0])
    image = x[0]

    gt = emojis[np.argmax(y[0])]['name']
    pred_class_name = emojis[pred_class]['name']

    image = plot_bounding_box(image, gt_coords, pred_coords, norm=True)

    # set text colors of labels
    color = 'green' if gt == pred_class_name else 'red'

    plt.imshow(image)
    plt.xlabel(f'Pred: {pred_class_name}', color=color)
    plt.ylabel(f'GT: {gt}', color=color)
    plt.xticks([])
    plt.yticks([])

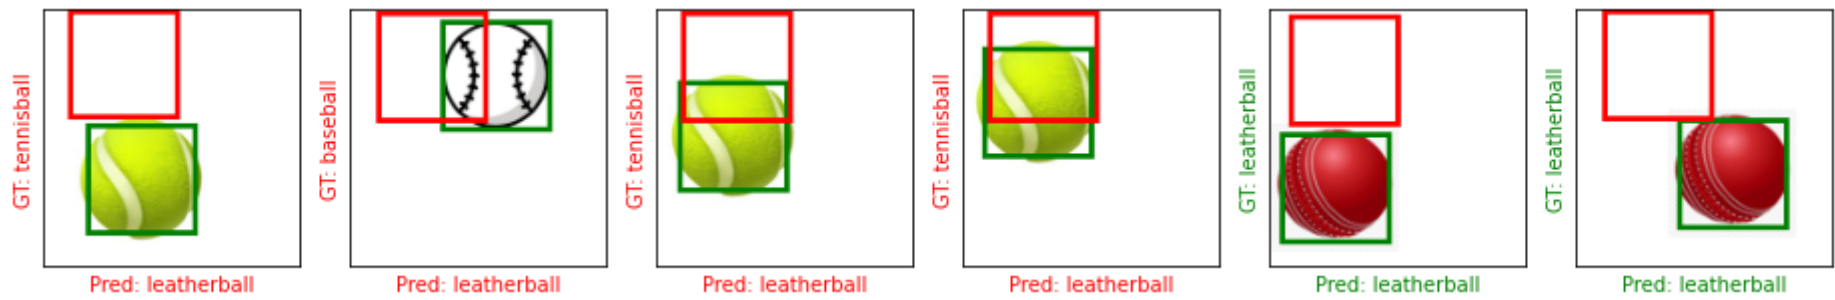
def test(model):
    test_datagen = data_generator(1) #1 = batch size is 1

    plt.figure(figsize=(16, 4))

```

```
# plot 6 images
for i in range(0, 6):
    plt.subplot(1, 6, i + 1)
    test_model(model, test_datagen)
plt.show()
```

```
test(model)
```



```
# create a custom call back
class ShowTestImages(tf.keras.callbacks.Callback): # to customize Callback class
    def on_epoch_end(self, epoch, logs=None): # on epoch end during the training, run test funtion
        test(self.model)
```

▼ Model Training

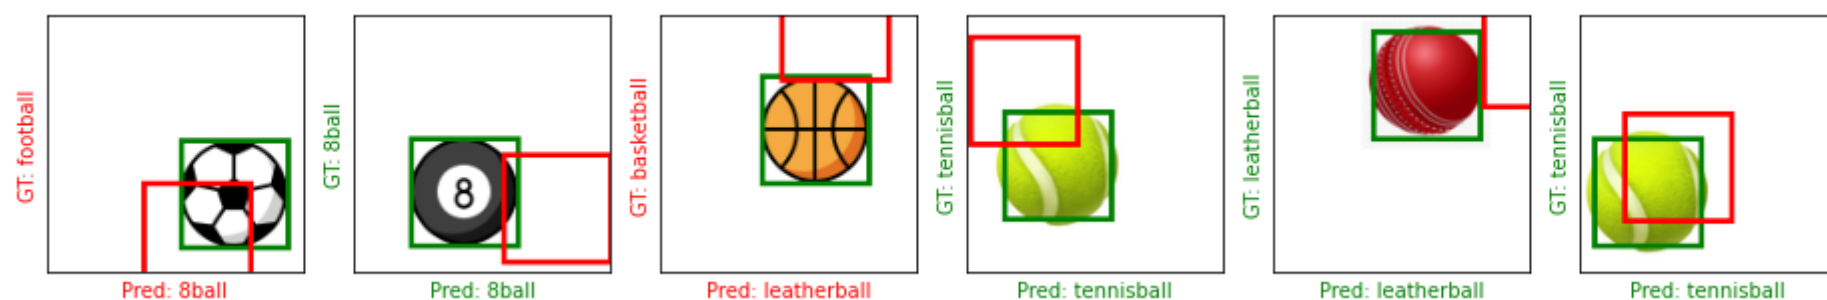
```
def lr_schedule(epoch, lr):
    if (epoch+1)%5 == 0:
        lr *= 0.2
    return max(lr, 3e-7)

_ = model.fit(
    data_generator(),
    epochs = 100,
    steps_per_epoch = 500,
    callbacks=[
        ShowTestImages(),
        tf.keras.callbacks.EarlyStopping(monitor='box_out_iou', patience=3, mode='max'),
        #stop is iou values does not increase for 4 consecutive epochs

        tf.keras.callbacks.LearningRateScheduler(lr_schedule)
    ]
)
```

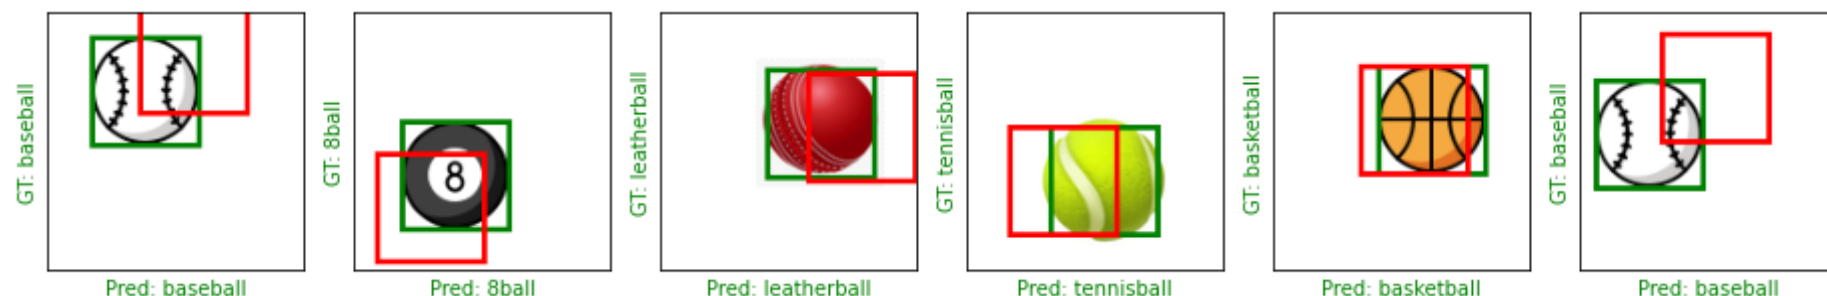

Epoch 1/100

500/500 [=====] - 14s 24ms/step - loss: 1.4389 - class_out_loss: 0.3602 - box_out_loss: 1.0787 - class_




Epoch 2/100

500/500 [=====] - 13s 25ms/step - loss: 0.0239 - class_out_loss: 0.0036 - box_out_loss: 0.0203 - class_




Epoch 3/100

500/500 [=====] - 12s 25ms/step - loss: 0.0133 - class_out_loss: 0.0036 - box_out_loss: 0.0097 - class_




Epoch 4/100

500/500 [=====] - 12s 24ms/step - loss: 0.0376 - class_out_loss: 0.0233 - box_out_loss: 0.0143 - class_



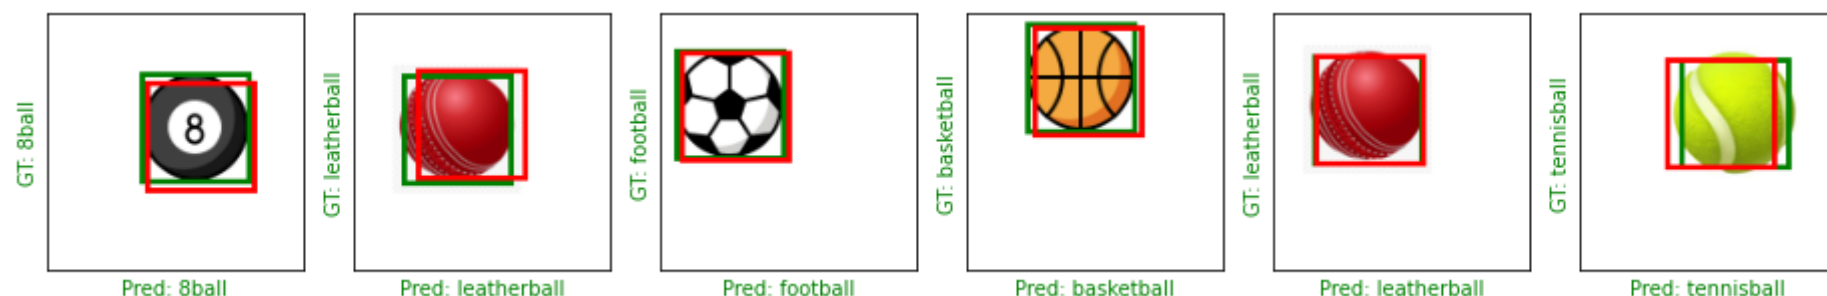
Epoch 5/100

500/500 [=====] - 13s 25ms/step - loss: 0.0049 - class_out_loss: 8.1411e-04 - box_out_loss: 0.0041 - cl




Epoch 6/100

500/500 [=====] - 12s 25ms/step - loss: 0.0026 - class_out_loss: 3.2887e-04 - box_out_loss: 0.0023 - cl



Epoch 7/100

500/500 [=====] - 13s 25ms/step - loss: 0.0236 - class_out_loss: 0.0210 - box_out_loss: 0.0026 - class_



Epoch 8/100

500/500 [=====] - 13s 25ms/step - loss: 0.0021 - class_out_loss: 3.7915e-04 - box_out_loss: 0.0017 - cl

