

## ▼ Download and Visualize Data

```
# import dataset

# !wget https://github.com/hfg-gmuend/openmoji/releases/latest/download/openmoji-72x72-color.zip
# !mkdir emojis
# !unzip -q openmoji-72x72-color.zip -d ./emojis
!pip install tensorflow==2.4

Requirement already satisfied: tensorboard-plugin-wit>=1.6.0 in /usr/local/lib/python3.7/dist-packages (from tensorboard~=2.4)
Requirement already satisfied: google-auth-oauthlib<0.5,>=0.4.1 in /usr/local/lib/python3.7/dist-packages (from tensorboard~=2.4)
Requirement already satisfied: setuptools>=41.0.0 in /usr/local/lib/python3.7/dist-packages (from tensorboard~=2.4->tensorflow==2.4)
Requirement already satisfied: markdown>=2.6.8 in /usr/local/lib/python3.7/dist-packages (from tensorboard~=2.4->tensorflow==2.4)
Requirement already satisfied: tensorboard-data-server<0.7.0,>=0.6.0 in /usr/local/lib/python3.7/dist-packages (from tensorboard~=2.4)
Requirement already satisfied: cachetools<5.0,>=2.0.0 in /usr/local/lib/python3.7/dist-packages (from google-auth<3,>=1.6.3->google-auth-oauthlib<0.5,>=0.4.1)
Requirement already satisfied: pyasn1-modules>=0.2.1 in /usr/local/lib/python3.7/dist-packages (from google-auth<3,>=1.6.3->google-auth-oauthlib<0.5,>=0.4.1)
Requirement already satisfied: rsa<5,>=3.1.4 in /usr/local/lib/python3.7/dist-packages (from google-auth<3,>=1.6.3->google-auth-oauthlib<0.5,>=0.4.1)
Requirement already satisfied: requests-oauthlib>=0.7.0 in /usr/local/lib/python3.7/dist-packages (from google-auth-oauthlib<0.5,>=0.4.1)
Requirement already satisfied: importlib-metadata>=4.4 in /usr/local/lib/python3.7/dist-packages (from markdown>=2.6.8->tensorflow==2.4)
Requirement already satisfied: zipp>=0.5 in /usr/local/lib/python3.7/dist-packages (from importlib-metadata>=4.4->markdown>=2.6.8->tensorflow==2.4)
Requirement already satisfied: pyasn1<0.5.0,>=0.4.6 in /usr/local/lib/python3.7/dist-packages (from pyasn1-modules>=0.2.1->google-auth-oauthlib<0.5,>=0.4.1)
Requirement already satisfied: idna<3,>=2.5 in /usr/local/lib/python3.7/dist-packages (from requests<3,>=2.21.0->tensorboard~=2.4)
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.7/dist-packages (from requests<3,>=2.21.0->tensorboard~=2.4)
Requirement already satisfied: chardet<4,>=3.0.2 in /usr/local/lib/python3.7/dist-packages (from requests<3,>=2.21.0->tensorboard~=2.4)
Requirement already satisfied: urllib3!=1.25.0,!1.25.1,<1.26,>=1.21.1 in /usr/local/lib/python3.7/dist-packages (from requests<3,>=2.21.0->tensorboard~=2.4)
Requirement already satisfied: oauthlib>=3.0.0 in /usr/local/lib/python3.7/dist-packages (from requests-oauthlib>=0.7.0->google-auth-oauthlib<0.5,>=0.4.1)
Building wheels for collected packages: wrapt
  Building wheel for wrapt (setup.py) ... done
  Created wheel for wrapt: filename=wrapt-1.12.1-cp37-cp37m-linux_x86_64.whl size=68717 sha256=6f5317f55e5aa751040c44da605926
  Stored in directory: /root/.cache/pip/wheels/62/76/4c/aa25851149f3f6d9785f6c869387ad82b3fd37582fa8147ac6
Successfully built wrapt
Installing collected packages: typing-extensions, grpcio, wrapt, tensorflow-estimator, h5py, gast, flatbuffers, tensorflow
  Attempting uninstall: typing-extensions
    Found existing installation: typing-extensions 3.10.0.2
    Uninstalling typing-extensions-3.10.0.2:
      Successfully uninstalled typing-extensions-3.10.0.2
  Attempting uninstall: grpcio
    Found existing installation: grpcio 1.42.0
    Uninstalling grpcio-1.42.0:
      Successfully uninstalled grpcio-1.42.0
  Attempting uninstall: wrapt
    Found existing installation: wrapt 1.13.3
    Uninstalling wrapt-1.13.3:
      Successfully uninstalled wrapt-1.13.3
  Attempting uninstall: tensorflow-estimator
    Found existing installation: tensorflow-estimator 2.7.0
    Uninstalling tensorflow-estimator-2.7.0:
      Successfully uninstalled tensorflow-estimator-2.7.0
  Attempting uninstall: h5py
    Found existing installation: h5py 3.1.0
    Uninstalling h5py-3.1.0:
      Successfully uninstalled h5py-3.1.0
  Attempting uninstall: gast
    Found existing installation: gast 0.4.0
    Uninstalling gast-0.4.0:
      Successfully uninstalled gast-0.4.0
  Attempting uninstall: flatbuffers
    Found existing installation: flatbuffers 2.0
    Uninstalling flatbuffers-2.0:
      Successfully uninstalled flatbuffers-2.0
  Attempting uninstall: tensorflow
    Found existing installation: tensorflow 2.7.0
    Uninstalling tensorflow-2.7.0:
      Successfully uninstalled tensorflow-2.7.0
Successfully installed flatbuffers-1.12 gast-0.3.3 grpcio-1.32.0 h5py-2.10.0 tensorflow-2.4.0 tensorflow-estimator-2.4.0 typi
```

```
from google.colab import files
```

```
uploaded = files.upload()
```

```
for fn in uploaded.keys():
    print('User uploaded file "{name}" with length {length} bytes'.format(
        name=fn, length=len(uploaded[fn])))
```

Choose Files No file chosen Upload widget is only available when the cell has been executed in the current browser session. Please rerun this cell to enable.  
 Saving 8BALL.png to 8BALL (1).png  
 Saving BASEBALL.png to BASEBALL (1).png  
 Saving BASKETBALL.png to BASKETBALL (1).png  
 Saving FOOTBALL.png to FOOTBALL (1).png  
 Saving LEATHERBALL.png to LEATHERBALL (1).png  
 Saving TENNISBALL.png to TENNISBALL (1).png  
 User uploaded file "8BALL.png" with length 2222 bytes

```
# import libraries
%matplotlib inline

import tensorflow as tf
import numpy as np
import matplotlib.pyplot as plt
import os

from PIL import Image, ImageDraw
from tensorflow.keras.layers import Input, Dense, Flatten, Conv2D, MaxPool2D, BatchNormalization, Dropout


print('Check if we are using TensorFlow 2.4')
print('Using TensorFlow version', tf.__version__)

    Check if we are using TensorFlow 2.4
    Using TensorFlow version 2.4.0

# balls actually using in this project
balls = {
    0: {'name': '8ball', 'file': '8BALL.png'},
    1: {'name': 'tennisball', 'file': 'TENNISBALL.png'},
    2: {'name': 'basketball', 'file': 'BASKETBALL.png'},
    3: {'name': 'football', 'file': 'FOOTBALL.png'},
    4: {'name': 'baseball', 'file': 'BASEBALL.png'},
    5: {'name': 'leatherball', 'file': 'LEATHERBALL.png'}
}

# place images in larger images - help synthesis data for localization
plt.figure(figsize=(9, 9))

for i, (j, e) in enumerate(balls.items()):
    plt.subplot(3, 3, i + 1)
    plt.imshow(plt.imread(e['file']))
    plt.xlabel(e['name'])
    plt.xticks([])
    plt.yticks([])
plt.show()
```



## ▼ Create Examples

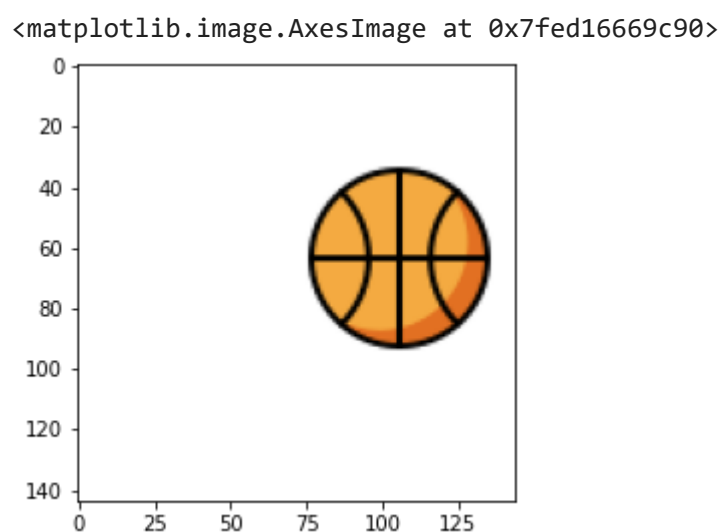
```
# loads balls and assign a key - image for each class in balls dictionary
for class_id, values in balls.items():
    png_file = Image.open(values['file']).convert('RGBA')
    png_file.load()
    new_file = Image.new("RGB", png_file.size, (255, 255, 255))
    new_file.paste(png_file, mask=png_file.split()[3])
    balls[class_id]['image'] = new_file
```

balls

```
{0: {'file': '8BALL.png',
      'image': <PIL.Image.Image image mode=RGB size=72x72 at 0x7FED1650D2D0>,
      'name': '8ball'},
 1: {'file': 'TENNISBALL.png',
      'image': <PIL.Image.Image image mode=RGB size=72x72 at 0x7FED1650D150>,
      'name': 'tennisball'},
 2: {'file': 'BASKETBALL.png',
      'image': <PIL.Image.Image image mode=RGB size=72x72 at 0x7FEDF031F5D0>,
      'name': 'basketball'},
 3: {'file': 'FOOTBALL.png',
      'image': <PIL.Image.Image image mode=RGB size=72x72 at 0x7FED1653A490>,
      'name': 'football'},
 4: {'file': 'BASEBALL.png',
      'image': <PIL.Image.Image image mode=RGB size=72x72 at 0x7FED1653A4D0>,
      'name': 'baseball'},
 5: {'file': 'LEATHERBALL.png',
      'image': <PIL.Image.Image image mode=RGB size=72x72 at 0x7FED1653A410>,
      'name': 'leatherball'}}
```

```
def create_example():
    class_id = np.random.randint(0,6) # randomly choose a emoji
    image = np.ones((144,144,3)) * 255 # create a white image
    row = np.random.randint(0,72) # place image randomly
    col = np.random.randint(0,72) #place image ramdomly
    # place emoji in blank emoji
    image[row: row+72, col: col+72, :] = np.array(balls[class_id]['image'])
    # return synthesize image
    return image.astype('uint8'), class_id, (row+5)/144 , (col+5)/144      # +10 becuae there is a white space of around 10 pixels in
    # /144 is normalization of image
```

```
image, class_id, row, col = create_example()
plt.imshow(image)
```



## ▼ Plot Bounding Boxes

```
#plotting bounding boxes
def plot_bounding_box(image, gt_coords, pred_coords=[], norm=False): # pass image, ground truth row col coordinates, predicted coordi
    if norm:
        image *=255 # if norm is true, we will denormalize it
        image=image.astype('uint8')
        image = Image.fromarray(image) #to convert image array to pil image
        draw = ImageDraw.Draw(image)

        # extrating row and col from groud truth values
        row, col = gt_coords

        # denormalizing coords
        row *= 144
        col *= 144
        draw.rectangle((col, row, col+62, row+62), outline = 'green', width=3) #+52 becuae iage is of 72 pixels becuae image has buffer


        # now same for pred coords
        if len(pred_coords)==2:
            # extrating row and col from groud truth values
            row, col = pred_coords

            # denormalizing coords
            row *= 144
            col *= 144
```

```
draw.rectangle((col, row, col+62, row+62), outline = 'red', width=3)
```

```
return image
```

```
image = plot_bounding_box(image, gt_coords = [row, col])
plt.imshow(image)
plt.title(balls[class_id]['name'])
plt.show()
```




## ▼ Data Generator

```
# create endless stream of these randomly generated eg which we will use in our model
def data_generator(batch_size=16):
    #run in endless loop and create example and labels of batch size
    while True:
        x_batch = np.zeros((batch_size, 144, 144, 3)) # 144 = size of image
        y_batch = np.zeros((batch_size, 6)) #9= no of class ids
        bbox_batch = np.zeros((batch_size, 2)) #2 = for row and col values

        # create examples of no of batch size
        for i in range(0, batch_size):
            image, class_id, row, col = create_example()
            x_batch[i] = image/255 # normalize image and 255 because they are pixel values
            y_batch[i, class_id] = 1.0
            bbox_batch[i] = np.array([row,col])
        yield {'image':x_batch}, {'class_out': y_batch, 'box_out': bbox_batch}
```

```
example, label = next(data_generator(1))
image = example['image'][0]
class_id = np.argmax(label['class_out'][0]) # to get the actual class id
coords = label['box_out'][0]
```

```
image = plot_bounding_box(image, coords, norm=True)
plt.imshow(image)
plt.title(balls[class_id]['name'])
plt.show()
```



## ▼ Model

```
# CNN model
input_ = Input(shape=(144,144,3), name='image')
```

```

x = input_

# we will have totoal 5 convolutional blocks
for i in range(0,4):
    n_filters = 2**(4+i)
    x = Conv2D(n_filters, 3, activation='relu')(x) # (x) - input is x
    x = BatchNormalization()(x)
    x = MaxPool2D(2)(x) #pool size od 2x2

x = Flatten()(x)
x = Dense(256, activation = 'relu')(x)

# now connect fully conected layer to our 2 output
class_out = Dense(6, activation='softmax', name = 'class_out')(x) # 9= outputs as we have 9 clases# for classifaction out, we use
box_out = Dense(2, name = 'box_out')(x) # we dont specify any activation as it is regression output and it is linear by default

# now construct the model
model = tf.keras.models.Model(input_, [class_out, box_out])
model.summary()

```

Model: "model\_4"

Layer (type)	Output Shape	Param #	Connected to
image (InputLayer)	[(None, 144, 144, 3)]	0	
conv2d_21 (Conv2D)	(None, 142, 142, 16)	448	image[0][0]
batch_normalization_20 (BatchNo	(None, 142, 142, 16)	64	conv2d_21[0][0]
max_pooling2d_20 (MaxPooling2D)	(None, 71, 71, 16)	0	batch_normalization_20[0][0]
conv2d_22 (Conv2D)	(None, 69, 69, 32)	4640	max_pooling2d_20[0][0]
batch_normalization_21 (BatchNo	(None, 69, 69, 32)	128	conv2d_22[0][0]
max_pooling2d_21 (MaxPooling2D)	(None, 34, 34, 32)	0	batch_normalization_21[0][0]
conv2d_23 (Conv2D)	(None, 32, 32, 64)	18496	max_pooling2d_21[0][0]
batch_normalization_22 (BatchNo	(None, 32, 32, 64)	256	conv2d_23[0][0]
max_pooling2d_22 (MaxPooling2D)	(None, 16, 16, 64)	0	batch_normalization_22[0][0]
conv2d_24 (Conv2D)	(None, 14, 14, 128)	73856	max_pooling2d_22[0][0]
batch_normalization_23 (BatchNo	(None, 14, 14, 128)	512	conv2d_24[0][0]
max_pooling2d_23 (MaxPooling2D)	(None, 7, 7, 128)	0	batch_normalization_23[0][0]
flatten_4 (Flatten)	(None, 6272)	0	max_pooling2d_23[0][0]
dense_4 (Dense)	(None, 256)	1605888	flatten_4[0][0]
class_out (Dense)	(None, 6)	1542	dense_4[0][0]
box_out (Dense)	(None, 2)	514	dense_4[0][0]
Total params: 1,706,344			
Trainable params: 1,705,864			
Non-trainable params: 480			

## ▼ Custom Metric: IoU

```

# intersection over union is the evaluation metric
# to measure the performance of the model - common in finding accuracy in object detector and object localizers
# iou is area of overlap [intesection of 2 boxes] between the predicted bounding box and actual values and
# combining the areas of both minus intersection will give area of union
# divide area of overlap by area of union - IoU values, if 1 - prediction is accurate

```

```

class IoU(tf.keras.metrics.Metric):
    def __init__(self, **kwargs):
        super(IoU, self).__init__(**kwargs)

        self.iou = self.add_weight(name='iou', initializer='zeros')
        self.total_iou = self.add_weight(name='total_iou', initializer='zeros')

```

```

self.num_ex = self.add_weight(name='num_ex', initializer='zeros')

def update_state(self, y_true, y_pred, sample_weight=None):
    def get_box(y):
        rows, cols = y[:, 0], y[:, 1]
        rows, cols = rows * 144, cols * 144
        y1, y2 = rows, rows + 62
        x1, x2 = cols, cols + 62
        return x1, y1, x2, y2

    def get_area(x1, y1, x2, y2):
        return tf.math.abs(x2 - x1) * tf.math.abs(y2 - y1)

    gt_x1, gt_y1, gt_x2, gt_y2 = get_box(y_true)
    p_x1, p_y1, p_x2, p_y2 = get_box(y_pred)

    i_x1 = tf.maximum(gt_x1, p_x1)
    i_y1 = tf.maximum(gt_y1, p_y1)
    i_x2 = tf.minimum(gt_x2, p_x2)
    i_y2 = tf.minimum(gt_y2, p_y2)

    i_area = get_area(i_x1, i_y1, i_x2, i_y2)
    u_area = get_area(gt_x1, gt_y1, gt_x2, gt_y2) + get_area(p_x1, p_y1, p_x2, p_y2) - i_area

    iou = tf.math.divide(i_area, u_area)
    self.num_ex.assign_add(1)
    self.total_iou.assign_add(tf.reduce_mean(iou))
    self.iou = tf.math.divide(self.total_iou, self.num_ex)

def result(self):
    return self.iou

def reset_state(self):
    self.iou = self.add_weight(name='iou', initializer='zeros')
    self.total_iou = self.add_weight(name='total_iou', initializer='zeros')
    self.num_ex = self.add_weight(name='num_ex', initializer='zeros')

```

## ▼ Task 8: Compile the Model

```

model.compile(
    # specify the loss for different outputs
    loss={
        'class_out': 'categorical_crossentropy', # or classification output
        'box_out': 'mse' #for regression output
    },
    optimizer = tf.keras.optimizers.Adam(learning_rate=1e-3),

    #set different metric for different output
    metrics={
        'class_out': 'accuracy',
        'box_out': IoU(name='iou')
    }
)

```

## ▼ Custom Callback: Model Testing

```

def test_model(model, test_datagen):
    example, label = next(test_datagen)
    x = example['image']
    y = label['class_out']
    box = label['box_out']

    pred_y, pred_box = model.predict(x)

    pred_coords = pred_box[0] # [0] because we want only one example
    gt_coords = box[0]
    pred_class = np.argmax(pred_y[0])
    image = x[0]

    gt = balls[np.argmax(y[0])]['name']
    pred_class_name = balls[pred_class]['name']

    image = plot_bounding_box(image, gt_coords, pred_coords, norm=True)

```

```

# set text colors of labels
color = 'green' if gt == pred_class_name else 'red'

plt.imshow(image)
plt.xlabel(f'Pred: {pred_class_name}', color=color)
plt.ylabel(f'GT: {gt}', color=color)
plt.xticks([])
plt.yticks([])

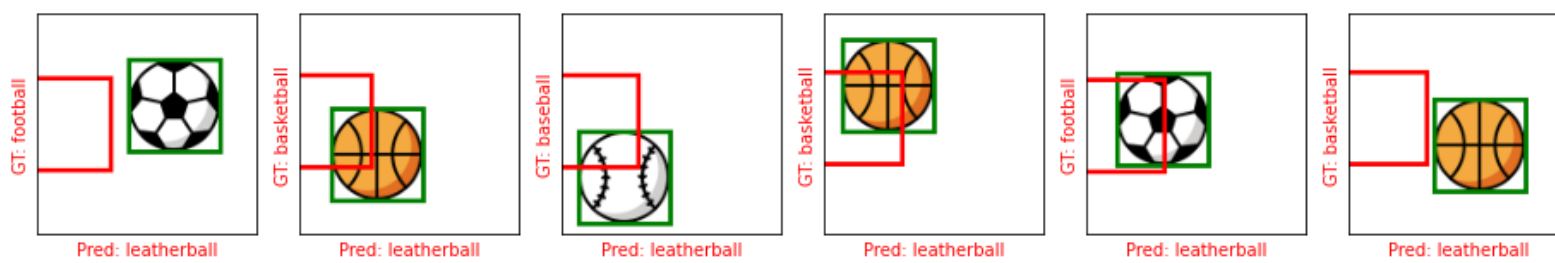
def test(model):
    test_datagen = data_generator(1) #1 = batch size is 1

    plt.figure(figsize=(16, 4))

    # plot 6 images
    for i in range(0, 6):
        plt.subplot(1, 6, i + 1)
        test_model(model, test_datagen)
    plt.show()

test(model)

```



```

# create a custom call back
class ShowTestImages(tf.keras.callbacks.Callback): # to customize Callback class
    def on_epoch_end(self, epoch, logs=None): # on epoch end during the training, run test funtion
        test(self.model)

```

## ▼ Model Training

```

def lr_schedule(epoch, lr):
    if (epoch+1)%5 == 0:
        lr *= 0.2
    return max(lr, 3e-7)

_ = model.fit(
    data_generator(),
    epochs = 100,
    steps_per_epoch = 500,
    callbacks=[
        ShowTestImages(),
        tf.keras.callbacks.EarlyStopping(monitor='box_out_iou', patience=3, mode='max'),
        #stop is iou values does not increase for 4 consecutive epochs


        tf.keras.callbacks.LearningRateScheduler(lr_schedule)
    ]
)

```



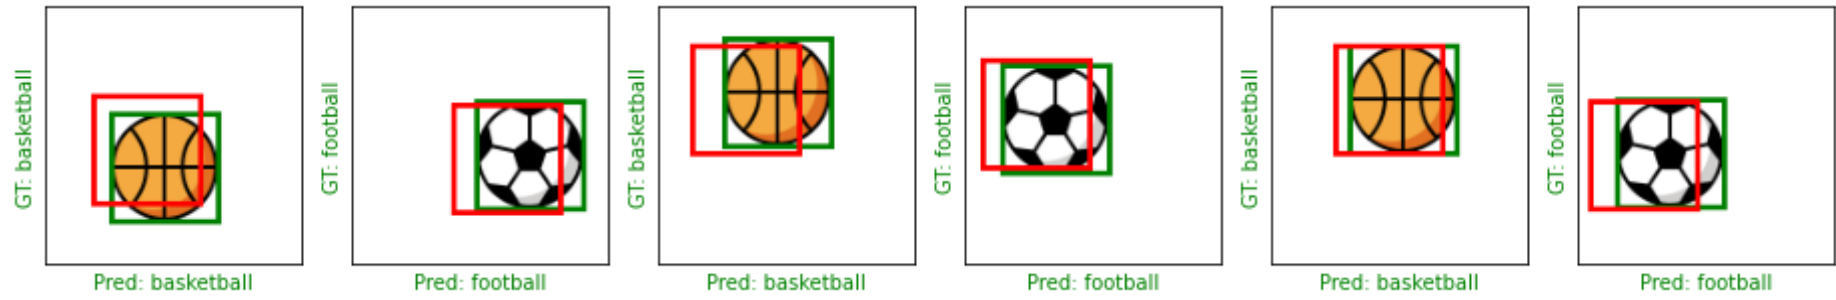
Epoch 1/100

500/500 [=====] - 13s 24ms/step - loss: 5.7395 - class\_out\_loss: 0.6591 - box\_out\_loss: 5.0804 - class\_




Epoch 2/100

500/500 [=====] - 12s 24ms/step - loss: 0.0088 - class\_out\_loss: 0.0017 - box\_out\_loss: 0.0072 - class\_




Epoch 3/100

500/500 [=====] - 12s 25ms/step - loss: 0.0035 - class\_out\_loss: 6.3001e-04 - box\_out\_loss: 0.0029 - cl



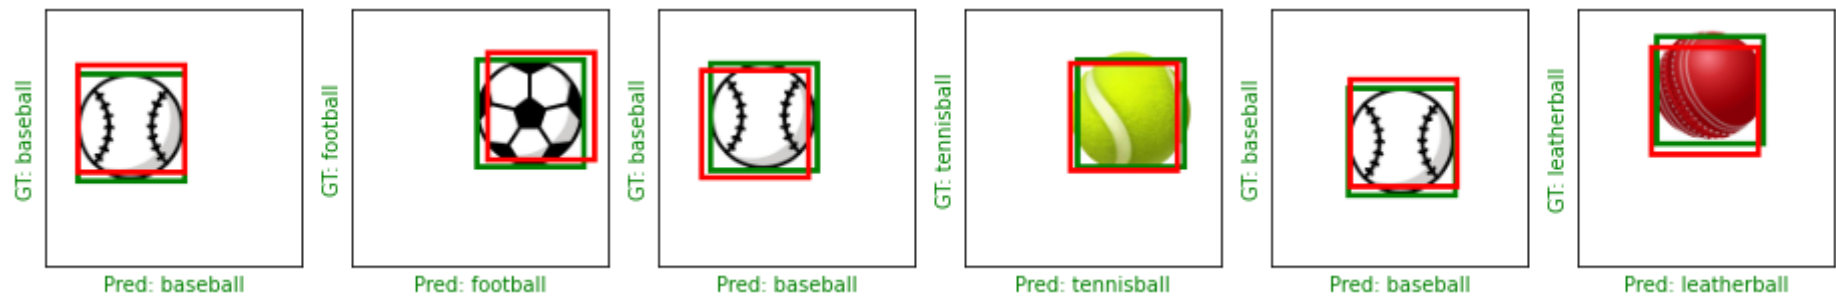
Epoch 4/100

500/500 [=====] - 12s 25ms/step - loss: 0.0019 - class\_out\_loss: 3.1236e-04 - box\_out\_loss: 0.0016 - cl



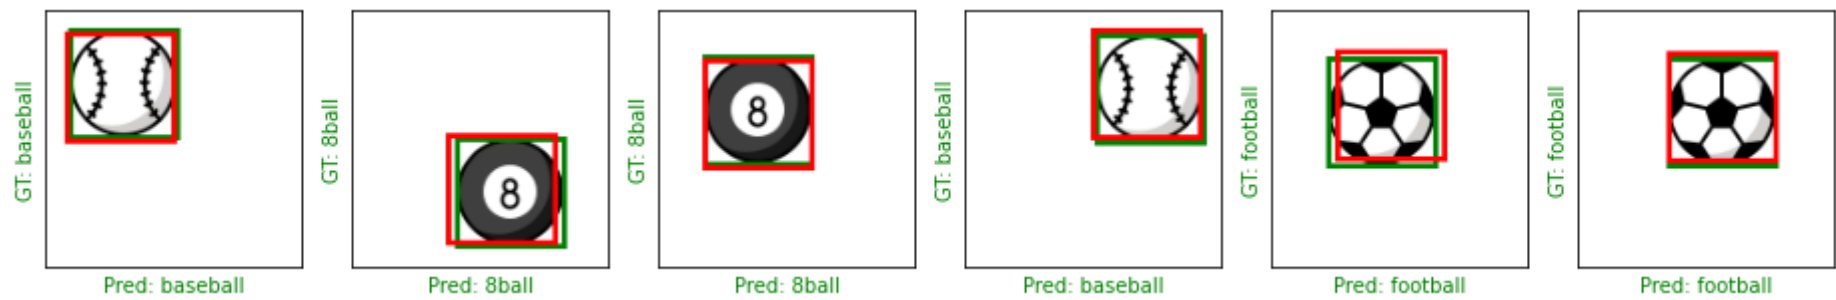
Epoch 5/100

500/500 [=====] - 12s 25ms/step - loss: 9.1429e-04 - class\_out\_loss: 2.2593e-04 - box\_out\_loss: 6.8836e



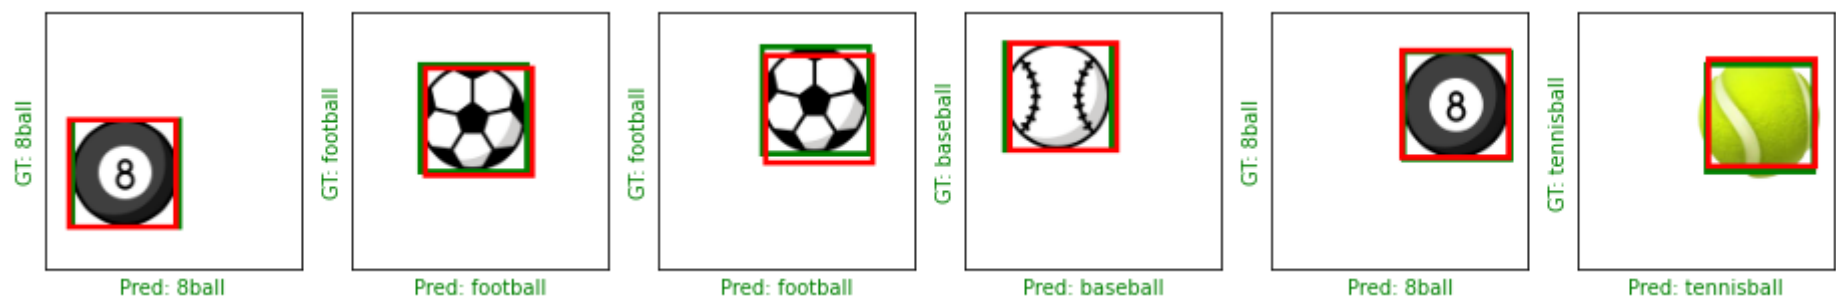
Epoch 6/100

500/500 [=====] - 12s 25ms/step - loss: 7.4677e-04 - class\_out\_loss: 1.4700e-04 - box\_out\_loss: 5.9976e




Epoch 7/100

500/500 [=====] - 12s 25ms/step - loss: 0.0020 - class\_out\_loss: 0.0013 - box\_out\_loss: 6.5552e-04 - cl



Epoch 8/100

500/500 [=====] - 13s 26ms/step - loss: 7.0534e-04 - class\_out\_loss: 1.8393e-04 - box\_out\_loss: 5.2141e





Pred: Baseball

Pred: Baseball

Pred: 8ball


Pred: Baseball

Pred: tennisball

Pred: football


Epoch 9/100

500/500 [=====] - 13s 26ms/step - loss: 0.0017 - class\_out\_loss: 0.0011 - box\_out\_loss: 5.4991e-04 - cl



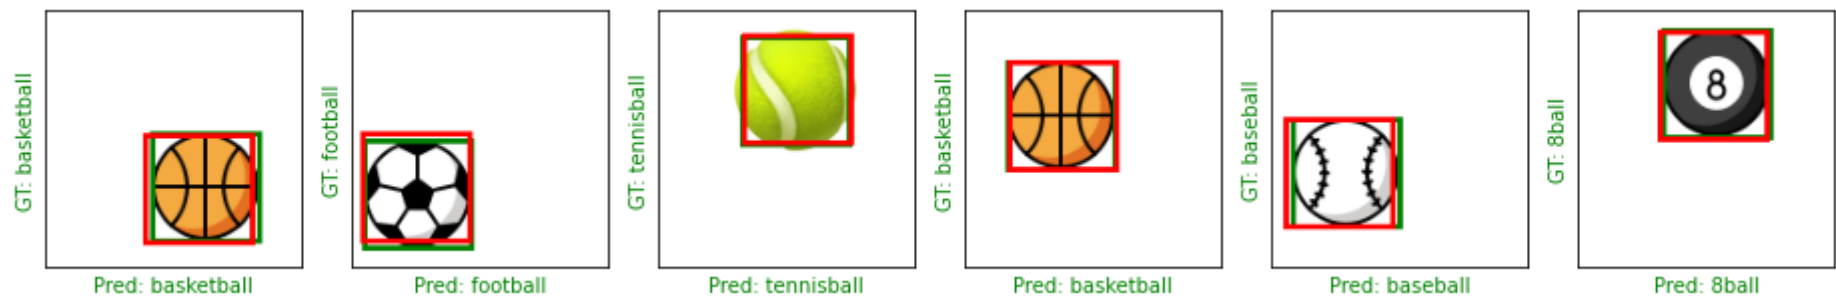
Epoch 10/100

500/500 [=====] - 13s 25ms/step - loss: 4.9440e-04 - class\_out\_loss: 9.7143e-05 - box\_out\_loss: 3.9726e



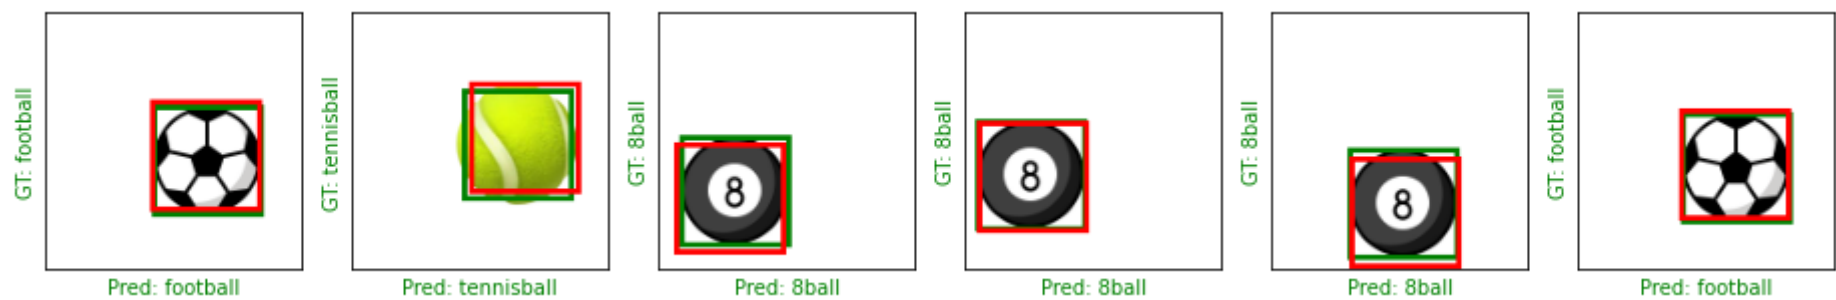
Epoch 11/100

500/500 [=====] - 13s 26ms/step - loss: 4.3102e-04 - class\_out\_loss: 8.0618e-05 - box\_out\_loss: 3.5041e




Epoch 12/100

500/500 [=====] - 13s 26ms/step - loss: 4.0492e-04 - class\_out\_loss: 7.0611e-05 - box\_out\_loss: 3.3431e



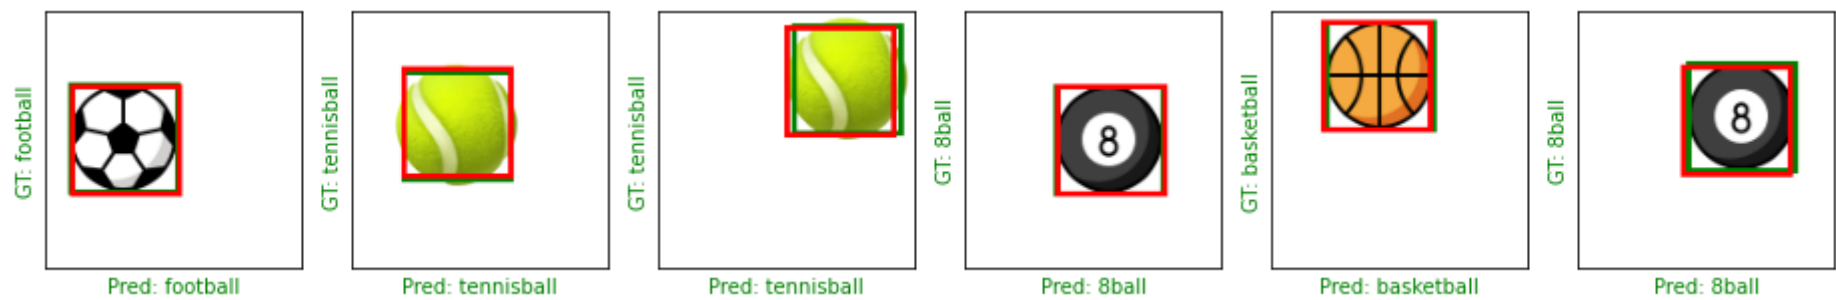
Epoch 13/100

500/500 [=====] - 13s 27ms/step - loss: 3.8628e-04 - class\_out\_loss: 7.5328e-05 - box\_out\_loss: 3.1095e




Epoch 14/100

500/500 [=====] - 13s 27ms/step - loss: 3.8424e-04 - class\_out\_loss: 7.3570e-05 - box\_out\_loss: 3.1067e




Epoch 15/100

500/500 [=====] - 14s 28ms/step - loss: 3.3730e-04 - class\_out\_loss: 5.9345e-05 - box\_out\_loss: 2.7796e




Epoch 16/100

500/500 [=====] - 15s 29ms/step - loss: 3.0895e-04 - class\_out\_loss: 5.2164e-05 - box\_out\_loss: 2.5679e

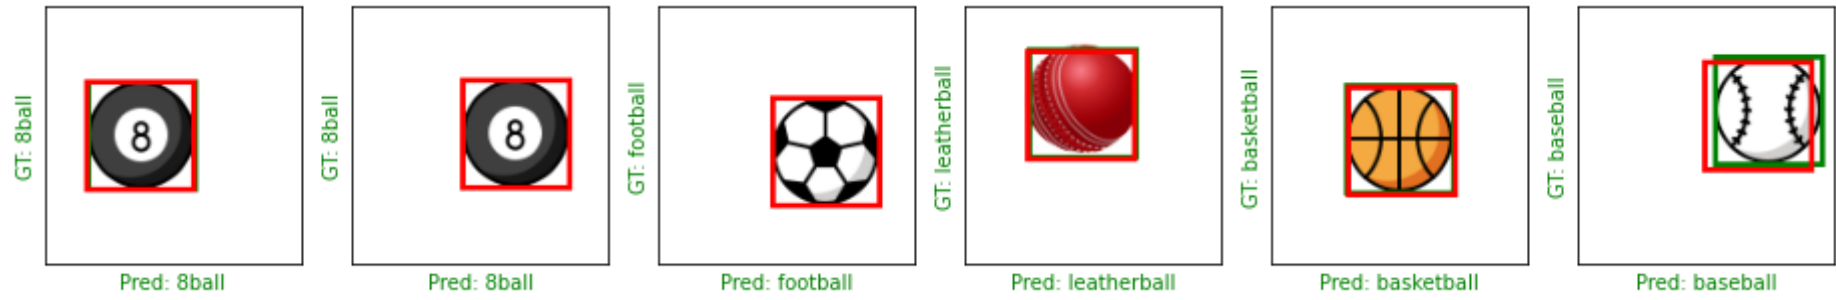





Epoch 25/100  
 500/500 [=====] - 16s 32ms/step - loss: 2.3737e-04 - class\_out\_loss: 3.0686e-05 - box\_out\_loss: 2.0668e




Epoch 26/100  
 500/500 [=====] - 16s 32ms/step - loss: 2.4427e-04 - class\_out\_loss: 3.5307e-05 - box\_out\_loss: 2.0896e



Epoch 27/100  
 500/500 [=====] - 16s 32ms/step - loss: 2.4596e-04 - class\_out\_loss: 3.3021e-05 - box\_out\_loss: 2.1294e



Epoch 28/100  
 500/500 [=====] - 16s 32ms/step - loss: 2.5350e-04 - class\_out\_loss: 4.0388e-05 - box\_out\_loss: 2.1311e



Epoch 29/100  
 500/500 [=====] - 16s 32ms/step - loss: 2.5304e-04 - class\_out\_loss: 4.2398e-05 - box\_out\_loss: 2.1064e