

▼ Download and Visualize Data

```
# import dataset

# !wget https://github.com/hfg-gmuend/openmoji/releases/latest/download/openmoji-72x72-color.zip
# !mkdir emojis
# !unzip -q openmoji-72x72-color.zip -d ./emojis
!pip install tensorflow==2.4

Requirement already satisfied: google-pasta~=0.2 in /usr/local/lib/python3.7/dist-packages (from tensorflow==2.4) (0.2.0)
Requirement already satisfied: protobuf>=3.9.2 in /usr/local/lib/python3.7/dist-packages (from tensorflow==2.4) (3.17.3)
Requirement already satisfied: tensorboard-plugin-wit>=1.6.0 in /usr/local/lib/python3.7/dist-packages (from tensorboard~=2.4->tensorflow==2.4) (1.6.0)
Requirement already satisfied: google-auth<3,>=1.6.3 in /usr/local/lib/python3.7/dist-packages (from tensorboard~=2.4->tensorflow==2.4) (1.29.0)
Requirement already satisfied: requests<3,>=2.21.0 in /usr/local/lib/python3.7/dist-packages (from tensorboard~=2.4->tensorflow==2.4) (2.27.1)
Requirement already satisfied: werkzeug>=0.11.15 in /usr/local/lib/python3.7/dist-packages (from tensorboard~=2.4->tensorflow==2.4) (0.16.1)
Requirement already satisfied: markdown>=2.6.8 in /usr/local/lib/python3.7/dist-packages (from tensorboard~=2.4->tensorflow==2.4) (3.3.7)
Requirement already satisfied: tensorboard-data-server<0.7.0,>=0.6.0 in /usr/local/lib/python3.7/dist-packages (from tensorboard~=2.4->tensorflow==2.4) (0.6.0)
Requirement already satisfied: google-auth-oauthlib<0.5,>=0.4.1 in /usr/local/lib/python3.7/dist-packages (from tensorboard~=2.4->tensorflow==2.4) (0.4.6)
Requirement already satisfied: setuptools>=41.0.0 in /usr/local/lib/python3.7/dist-packages (from tensorboard~=2.4->tensorflow==2.4) (57.5.0)
Requirement already satisfied: pyasn1-modules>=0.2.1 in /usr/local/lib/python3.7/dist-packages (from google-auth<3,>=1.6.3->tensorflow==2.4) (0.3.1)
Requirement already satisfied: cachetools<5.0,>=2.0.0 in /usr/local/lib/python3.7/dist-packages (from google-auth<3,>=1.6.3->tensorflow==2.4) (4.2.4)
Requirement already satisfied: rsa<5,>=3.1.4 in /usr/local/lib/python3.7/dist-packages (from google-auth<3,>=1.6.3->tensorflow==2.4) (4.9)
Requirement already satisfied: requests-oauthlib>=0.7.0 in /usr/local/lib/python3.7/dist-packages (from google-auth-oauthlib<0.5,>=0.4.1->tensorflow==2.4) (1.3.1)
Requirement already satisfied: importlib-metadata>=4.4 in /usr/local/lib/python3.7/dist-packages (from markdown>=2.6.8->tensorflow==2.4) (6.7.0)
Requirement already satisfied: zipp>=0.5 in /usr/local/lib/python3.7/dist-packages (from importlib-metadata>=4.4->markdown>=2.6.8->tensorflow==2.4) (3.15.0)
Requirement already satisfied: pyasn1<0.5.0,>=0.4.6 in /usr/local/lib/python3.7/dist-packages (from pyasn1-modules>=0.2.1->tensorflow==2.4) (0.4.8)
Requirement already satisfied: idna<3,>=2.5 in /usr/local/lib/python3.7/dist-packages (from requests<3,>=2.21.0->tensorflow==2.4) (3.4)
Requirement already satisfied: chardet<4,>=3.0.2 in /usr/local/lib/python3.7/dist-packages (from requests<3,>=2.21.0->tensorflow==2.4) (3.7.4)
Requirement already satisfied: urllib3!=1.25.0,!1.25.1,<1.26,>=1.21.1 in /usr/local/lib/python3.7/dist-packages (from requests<3,>=2.21.0->tensorflow==2.4) (1.26.13)
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.7/dist-packages (from requests<3,>=2.21.0->tensorflow==2.4) (2022.9.24)
Requirement already satisfied: oauthlib>=3.0.0 in /usr/local/lib/python3.7/dist-packages (from requests-oauthlib>=0.7.0->tensorflow==2.4) (3.2.2)
Building wheels for collected packages: wrapt
  Building wheel for wrapt (setup.py) ... done
  Created wheel for wrapt: filename=wrapt-1.12.1-cp37-cp37m-linux_x86_64.whl size=68717 sha256=69fdfe127412e9121b22eec09bac9a2b3f37582fa8147ac6
  Stored in directory: /root/.cache/pip/wheels/62/76/4c/aa25851149f3f6d9785f6c869387ad82b3fd37582fa8147ac6
Successfully built wrapt
Installing collected packages: typing-extensions, grpcio, wrapt, tensorflow-estimator, h5py, gast, flatbuffers, tensorflow
  Attempting uninstall: typing-extensions
    Found existing installation: typing-extensions 3.10.0.2
    Uninstalling typing-extensions-3.10.0.2:
      Successfully uninstalled typing-extensions-3.10.0.2
  Attempting uninstall: grpcio
    Found existing installation: grpcio 1.42.0
    Uninstalling grpcio-1.42.0:
      Successfully uninstalled grpcio-1.42.0
  Attempting uninstall: wrapt
    Found existing installation: wrapt 1.13.3
    Uninstalling wrapt-1.13.3:
      Successfully uninstalled wrapt-1.13.3
  Attempting uninstall: tensorflow-estimator
    Found existing installation: tensorflow-estimator 2.7.0
    Uninstalling tensorflow-estimator-2.7.0:
      Successfully uninstalled tensorflow-estimator-2.7.0
  Attempting uninstall: h5py
    Found existing installation: h5py 3.1.0
    Uninstalling h5py-3.1.0:
      Successfully uninstalled h5py-3.1.0
  Attempting uninstall: gast
    Found existing installation: gast 0.4.0
    Uninstalling gast-0.4.0:
      Successfully uninstalled gast-0.4.0
  Attempting uninstall: flatbuffers
    Found existing installation: flatbuffers 2.0
    Uninstalling flatbuffers-2.0:
      Successfully uninstalled flatbuffers-2.0
  Attempting uninstall: tensorflow
```

```
from google.colab import files
```

```
uploaded = files.upload()
```

```
for fn in uploaded.keys():
    print('User uploaded file "{name}" with length {length} bytes'.format(
        name=fn, length=len(uploaded[fn])))
```

Choose Files 6 files

- **8BALL.png**(image/png) - 2232 bytes, last modified: 4/30/2020 - 100% done
- **BASEBALL.png**(image/png) - 2387 bytes, last modified: 6/11/2020 - 100% done
- **BASKETBALL.png**(image/png) - 2132 bytes, last modified: 4/30/2020 - 100% done
- **FOOTBALL.png**(image/png) - 2193 bytes, last modified: 4/29/2020 - 100% done
- **LEATHERBALL.png**(image/png) - 7495 bytes, last modified: 12/12/2021 - 100% done
- **TENNISBALL.png**(image/png) - 9556 bytes, last modified: 12/12/2021 - 100% done

Saving 8BALL.png to 8BALL.png

Saving BASEBALL.png to BASEBALL.png

Saving BASKETBALL.png to BASKETBALL.png

Saving FOOTBALL.png to FOOTBALL.png

Saving LEATHERBALL.png to LEATHERBALL.png

Saving TENNISBALL.png to TENNISBALL.png

User uploaded file "8BALL.png" with length 2232 bytes

User uploaded file "BASEBALL.png" with length 2387 bytes

```
# import libraries
```

```
%matplotlib inline
```

```
import tensorflow as tf
```

```
import numpy as np
```

```
import matplotlib.pyplot as plt
```

```
import os
```

```
from PIL import Image, ImageDraw
```

```
from tensorflow.keras.layers import Input, Dense, Flatten, Conv2D, MaxPool2D, BatchNormalization, Dropout
```

```
print('Check if we are using TensorFlow 2.4')
```

```
print('Using TensorFlow version', tf.__version__)
```

```
Check if we are using TensorFlow 2.4
```

```
Using TensorFlow version 2.4.0
```

```
# balls actually using in this project
```

```
balls = {
```

```
    0: {'name': '8ball', 'file': '8BALL.png'},
```

```
    1: {'name': 'tennisball', 'file': 'TENNISBALL.png'},
```

```
    2: {'name': 'basketball', 'file': 'BASKETBALL.png'},
```

```
    3: {'name': 'football', 'file': 'FOOTBALL.png'},
```

```
    4: {'name': 'baseball', 'file': 'BASEBALL.png'},
```

```
    5: {'name': 'leatherball', 'file': 'LEATHERBALL.png'}
}
```

```
# place images in larger images - help synthesis data for localization
```

```
plt.figure(figsize=(9, 9))
```

```
for i, (j, e) in enumerate(balls.items()):
```

```
    plt.subplot(3, 3, i + 1)
```

```
    plt.imshow(plt.imread(e['file']))
```

```
    plt.xlabel(e['name'])
```

```
    plt.xticks([])
```

```
    plt.yticks([])
```

```
plt.show()
```



8ball




tennisball




basketball



football



baseball



leatherball

▼ Create Examples

```
# loads balls and assign a key - image for each class in balls dictionary
```

```
for class_id, values in balls.items():
```

```

png_file = Image.open(values['file']).convert('RGBA')
png_file.load()
new_file = Image.new("RGB", png_file.size, (255, 255, 255))
new_file.paste(png_file, mask=png_file.split()[3])
balls[class_id]['image'] = new_file

```

balls

```

{0: {'file': '8BALL.png',
     'image': <PIL.Image.Image image mode=RGB size=72x72 at 0x7FA7C2D9D950>,
     'name': '8ball'},
 1: {'file': 'TENNISBALL.png',
     'image': <PIL.Image.Image image mode=RGB size=72x72 at 0x7FA7C45EB910>,
     'name': 'tennisball'},
 2: {'file': 'BASKETBALL.png',
     'image': <PIL.Image.Image image mode=RGB size=72x72 at 0x7FA7C2D9D990>,
     'name': 'basketball'},
 3: {'file': 'FOOTBALL.png',
     'image': <PIL.Image.Image image mode=RGB size=72x72 at 0x7FA7C45C5410>,
     'name': 'football'},
 4: {'file': 'BASEBALL.png',
     'image': <PIL.Image.Image image mode=RGB size=72x72 at 0x7FA7C45D0690>,
     'name': 'baseball'},
 5: {'file': 'LEATHERBALL.png',
     'image': <PIL.Image.Image image mode=RGB size=72x72 at 0x7FA7C2DAAD90>,
     'name': 'leatherball'}}

```

```

def create_example():
    class_id = np.random.randint(0,6) # randomly choose a emoji
    image = np.ones((144,144,3)) * 255 # create a white image
    row = np.random.randint(0,72) # place image randomly
    col = np.random.randint(0,72) #place image ramdomly
    # place emoji in blank emoji
    image[row:row+72, col:col+72, :]= np.array(balls[class_id]['image'])
    # return synthesize image
    return image.astype('uint8'), class_id, (row+5)/144 , (col+5)/144      # +10 becuase there is a white space of around 10 pixels in
    # /144 is normalization of image


```

```

image, class_id, row, col = create_example()
plt.imshow(image)

```

<matplotlib.image.AxesImage at 0x7fa7c4633a50>



▼ Plot Bounding Boxes

```

#plotting bounding boxes
def plot_bounding_box(image, gt_coords, pred_coords=[], norm=False): # pass image, ground truth row col coordinates, predicted coordi
    if norm:
        image *=255 # if norm is true, we will denormalize it
        image=image.astype('uint8')
        image = Image.fromarray(image) #to convert image array to pil image
        draw = ImageDraw.Draw(image)

        # extrating row and col from groud truth values
        row, col = gt_coords

        # denormalizing coords
        row *= 144
        col *= 144
        draw.rectangle((col, row, col+62, row+62), outline = 'green', width=3) #+52 becuase iage is of 72 pixels becuase image has buffer

        # now same for pred coords

```

```

if len(pred_coords)==2:
    # extrating row and col from groud truth values
    row, col = pred_coords

    # denormalizing coords
    row *= 144
    col *= 144
    draw.rectangle((col, row, col+62, row+62), outline = 'red', width=3)


return image

```

```

image = plot_bounding_box(image, gt_coords = [row, col])
plt.imshow(image)
plt.title(balls[class_id]['name'])
plt.show()

```



▼ Data Generator

```

# create endless stream of these randomly generated eg which we will use in our model
def data_generator(batch_size=16):
    #run in endless loop and create example and labels of batch size
    while True:
        x_batch = np.zeros((batch_size, 144, 144, 3)) # 144 = size of image
        y_batch = np.zeros((batch_size, 6)) #9= no of class ids
        bbox_batch = np.zeros((batch_size, 2)) #2 = for row and col values

        # create examples of no of batch size
        for i in range(0, batch_size):
            image, class_id, row, col = create_example()
            x_batch[i] = image/255 # normalize image and 255 because they are pixel values
            y_batch[i, class_id] = 1.0
            bbox_batch[i] = np.array([row,col])
        yield {'image':x_batch} ,{'class_out': y_batch, 'box_out': bbox_batch}

```

```


example, label = next(data_generator(1))
image = example['image'][0]
class_id = np.argmax(label['class_out'][0]) # to get the actual class id
coords = label['box_out'][0]

```

```

image = plot_bounding_box(image, coords, norm=True)
plt.imshow(image)
plt.title(balls[class_id]['name'])
plt.show()

```



▼ Model

```
# CNN model
input_ = Input(shape=(144,144,3), name='image')

x = input_

# we will have totoal 4 convolutional blocks
for i in range(0,4):
    n_filters = 2**(4+i)
    x = Conv2D(n_filters, 3, activation='relu')(x) # (x) - input is x
    x = BatchNormalization()(x)
    x = MaxPool2D(2)(x) #pool size od 2x2

x = Flatten()(x)
x = Dense(256, activation = 'relu')(x)

# now connect fully conected layer to our 2 output
class_out = Dense(6, activation='softmax', name = 'class_out')(x) # 9= outputs as we have 9 clases# for classifaction out, we use
box_out = Dense(2, name = 'box_out')(x) # we dont specify any activation as it is regression output and it is linear by default

# now construct the model
model = tf.keras.models.Model(input_, [class_out, box_out])
model.summary()
```

Model: "model"

Layer (type)	Output Shape	Param #	Connected to
=====			
image (InputLayer)	[(None, 144, 144, 3)]	0	
conv2d (Conv2D)	(None, 142, 142, 16)	448	image[0][0]
batch_normalization (BatchNormaliza	(None, 142, 142, 16)	64	conv2d[0][0]
max_pooling2d (MaxPooling2D)	(None, 71, 71, 16)	0	batch_normalization[0][0]
conv2d_1 (Conv2D)	(None, 69, 69, 32)	4640	max_pooling2d[0][0]
batch_normalization_1 (BatchNor	(None, 69, 69, 32)	128	conv2d_1[0][0]
max_pooling2d_1 (MaxPooling2D)	(None, 34, 34, 32)	0	batch_normalization_1[0][0]
conv2d_2 (Conv2D)	(None, 32, 32, 64)	18496	max_pooling2d_1[0][0]
batch_normalization_2 (BatchNor	(None, 32, 32, 64)	256	conv2d_2[0][0]
max_pooling2d_2 (MaxPooling2D)	(None, 16, 16, 64)	0	batch_normalization_2[0][0]
conv2d_3 (Conv2D)	(None, 14, 14, 128)	73856	max_pooling2d_2[0][0]
batch_normalization_3 (BatchNor	(None, 14, 14, 128)	512	conv2d_3[0][0]
max_pooling2d_3 (MaxPooling2D)	(None, 7, 7, 128)	0	batch_normalization_3[0][0]
flatten (Flatten)	(None, 6272)	0	max_pooling2d_3[0][0]
dense (Dense)	(None, 256)	1605888	flatten[0][0]
class_out (Dense)	(None, 6)	1542	dense[0][0]
box_out (Dense)	(None, 2)	514	dense[0][0]
=====			
Total params: 1,706,344			
Trainable params: 1,705,864			
Non-trainable params: 480			

▼ Custom Metric: IoU

```
# intersection over union is the evaluation metric
# to measure the performance of the model - common in finding accuracy in object detector and object localizers
# iou is area of overlap [intesection of 2 boxes] between the predicted bounding box and actual values and
```

```

# combining the areas of both minus intersection will give area of union
# divide area of overlap by area of union. Tell value if 1. Prediction is accurate

class IoU(tf.keras.metrics.Metric):
    def __init__(self, **kwargs):
        super(IoU, self).__init__(**kwargs)

        self.iou = self.add_weight(name='iou', initializer='zeros')
        self.total_iou = self.add_weight(name='total_iou', initializer='zeros')
        self.num_ex = self.add_weight(name='num_ex', initializer='zeros')

    def update_state(self, y_true, y_pred, sample_weight=None):
        def get_box(y):
            rows, cols = y[:, 0], y[:, 1]
            rows, cols = rows * 144, cols * 144
            y1, y2 = rows, rows + 62
            x1, x2 = cols, cols + 62
            return x1, y1, x2, y2

        def get_area(x1, y1, x2, y2):
            return tf.math.abs(x2 - x1) * tf.math.abs(y2 - y1)

        gt_x1, gt_y1, gt_x2, gt_y2 = get_box(y_true)
        p_x1, p_y1, p_x2, p_y2 = get_box(y_pred)

        i_x1 = tf.maximum(gt_x1, p_x1)
        i_y1 = tf.maximum(gt_y1, p_y1)
        i_x2 = tf.minimum(gt_x2, p_x2)
        i_y2 = tf.minimum(gt_y2, p_y2)

        i_area = get_area(i_x1, i_y1, i_x2, i_y2)
        u_area = get_area(gt_x1, gt_y1, gt_x2, gt_y2) + get_area(p_x1, p_y1, p_x2, p_y2) - i_area

        iou = tf.math.divide(i_area, u_area)
        self.num_ex.assign_add(1)
        self.total_iou.assign_add(tf.reduce_mean(iou))
        self.iou = tf.math.divide(self.total_iou, self.num_ex)

    def result(self):
        return self.iou

    def reset_state(self):
        self.iou = self.add_weight(name='iou', initializer='zeros')
        self.total_iou = self.add_weight(name='total_iou', initializer='zeros')
        self.num_ex = self.add_weight(name='num_ex', initializer='zeros')

```

▼ Task 8: Compile the Model

```

model.compile(
    # specify the loss for fiffereent outputs
    loss={
        'class_out': 'categorical_crossentropy', # or classifaction output
        'box_out': 'mse' #for regression output
    },
    optimizer = tf.keras.optimizers.Adam(learning_rate=1e-3),

    #set differnet metric for different output
    metrics={
        'class_out': 'accuracy',
        'box_out': IoU(name='iou')
    }
)

```

▼ Custom Callback: Model Testing

```

def test_model(model, test_datagen):
    example, label = next(test_datagen)
    x = example['image']
    y = label['class_out']
    box = label['box_out']

    pred_y, pred_box = model.predict(x)

    pred_coords = pred_box[0] # [0] beacuse we want only one example
    gt_coords = box[0]

```

```

pred_class = np.argmax(pred_y[0])
image = x[0]

gt = balls[np.argmax(y[0])]['name']
pred_class_name = balls[pred_class]['name']

image = plot_bounding_box(image, gt_coords, pred_coords, norm=True)

# set text colors of labels
color = 'green' if gt == pred_class_name else 'red'

plt.imshow(image)
plt.xlabel(f'Pred: {pred_class_name}', color=color)
plt.ylabel(f'GT: {gt}', color=color)
plt.xticks([])
plt.yticks([])

```

```

def test(model):
    test_datagen = data_generator(1) #1 = batch size is 1

```


```
plt.figure(figsize=(16, 4))
```

```

# plot 6 images
for i in range(0, 6):
    plt.subplot(1, 6, i + 1)
    test_model(model, test_datagen)
plt.show()

```

```
test(model)
```



```

# create a custom call back
class ShowTestImages(tf.keras.callbacks.Callback): # to customize Callback class
    def on_epoch_end(self, epoch, logs=None): # on epoch end during the training, run test funtion
        test(self.model)

```

▼ Model Training


```

def lr_schedule(epoch, lr):
    if (epoch+1)%5 == 0:
        lr *= 0.2
    return max(lr, 3e-7)

_ = model.fit(
    data_generator(),
    epochs = 100,
    steps_per_epoch = 500,
    callbacks=[
        ShowTestImages(),
        tf.keras.callbacks.EarlyStopping(monitor='box_out_iou', patience=3, mode='max'),
        #stop is iou values does not increase for 4 consecutive epochs


        tf.keras.callbacks.LearningRateScheduler(lr_schedule)
    ]
)

```




Epoch 24/100

500/500 [=====] - 14s 28ms/step - loss: 4.2631e-04 - class_out_loss: 2.7818e-05 - box_out_loss: 3.9849




Epoch 25/100

500/500 [=====] - 14s 28ms/step - loss: 4.3857e-04 - class_out_loss: 2.7475e-05 - box_out_loss: 4.1110



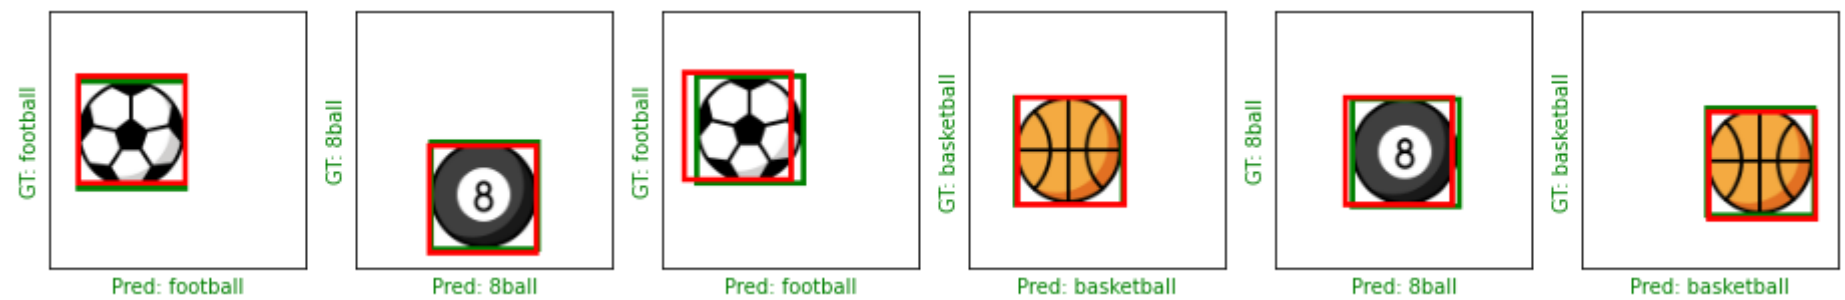
Epoch 26/100

500/500 [=====] - 14s 28ms/step - loss: 6.2530e-04 - class_out_loss: 2.5180e-04 - box_out_loss: 3.7350



Epoch 27/100

500/500 [=====] - 14s 28ms/step - loss: 4.0534e-04 - class_out_loss: 2.3729e-05 - box_out_loss: 3.8162



Epoch 28/100

500/500 [=====] - 14s 28ms/step - loss: 4.3162e-04 - class_out_loss: 4.7983e-05 - box_out_loss: 3.8364

